JOB DESCRIPTION
Data Scientist (GenAI)

Designation – Data Scientist (GenAI)
Experience – 2 + Years

Job Summary:

We are looking for experienced Data Scientists specializing in Generative AI to join our team. The ideal candidates should have 2+ years of experience in building and deploying Generative AI solutions. This role requires expertise in transformer architectures, LLM fine-tuning, cloud-based AI deployments, and MLOps for GenAI.

Job Description:

1. Develop, fine-tune, and deploy Generative AI models (GPT, BERT, T5, LLaMA, etc.).
2. Implement LLM fine-tuning, prompt engineering, and reinforcement learning techniques for domain-specific tasks.
3. Design, optimize, and scale AI/ML pipelines for large-scale data processing and model inference.
4. Work with Hugging Face Transformers, LangChain, and OpenAI APIs for building NLP and multimodal AI solutions.
5. Develop and maintain CI/CD pipelines for AI/ML models, ensuring efficient model versioning, monitoring, and deployment.
6. Utilize cloud platforms (Azure, AWS, or GCP) for scalable and secure AI model hosting.
7. Implement MLOps best practices for managing model lifecycle, monitoring performance, and retraining workflows.
8. Collaborate with data engineers to optimize data pipelines, ETL processes, and database queries for AI applications.
9. Research and implement state-of-the-art transformer architectures to improve model accuracy and efficiency.
10. Analyze and interpret structured, unstructured, and semi-structured data to drive AI-driven business insights.
11. Design and conduct experiments with statistical rigor, implementing hypothesis-driven solutions.

CELEBAL TECHNOLOGIES

Must Have:

1. Cloud Expertise (Mandatory): Strong experience in Azure/AWS/GCP for AI deployments.
2. Programming: Advanced Python skills with expertise in libraries like Pandas, NumPy, Scikit-learn, TensorFlow, PyTorch, Hugging Face.
3. Data Engineering: Proficiency in SQL, NoSQL, and big data tools (Spark, Databricks, or Snowflake).
4. Transformer Architectures: Deep understanding of GPT, BERT, T5, LLaMA, and diffusion models.
5. LLM Fine-tuning: Hands-on experience in fine-tuning and optimizing large language models.
6. DevOps for AI/ML: Experience with CI/CD, MLflow, model versioning, and AI model monitoring.
7. Analytical Skills: Strong grasp of statistics, probability, and machine learning algorithms.
8. Data Handling: Ability to process structured, unstructured, and semi-structured data efficiently.