JOB DESCRIPTION
Data Scientist / Machine Learning Engineer

Designation – Data Scientist/ Machine Learning Engineer
Experience – 2 to 5 Years

Job Summary:

We are looking for a highly skilled and motivated Data Scientist / Machine Learning Engineer with expertise in classical Machine Learning, Deep Learning, Time Series Forecasting, and Vision AI. The ideal candidate should have at least 2+ years of experience in Python programming, AI/ML model development, and cloud-based deployments.

As a key member of our team, you will be responsible for developing, optimizing, and deploying AI/ML solutions across multiple domains, including predictive analytics, time series forecasting, and computer vision. You will also work on building scalable APIs, deploying models on cloud platforms (Azure, AWS, or GCP), and implementing MLOps best practices to streamline AI integration into production environments.

Job Description:

1. Develop and deploy Machine Learning (ML) and Deep Learning (DL) models for predictive analytics, classification, regression, and clustering tasks.
2. Work on Time Series Forecasting using ARIMA, LSTMs, Prophet, and advanced forecasting techniques.
3. Build and optimize Computer Vision models for object detection, image classification, and scene understanding.
4. Implement Natural Language Processing (NLP) models for text analytics, OCR, and document intelligence applications.
5. Design, develop, and maintain high-performance Python APIs (FastAPI, Flask, Django) to expose AI models as scalable services.
6. Work with structured and unstructured data, performing feature engineering and data transformations.
7. Deploy AI models on Azure/AWS/GCP, ensuring cost optimization, scalability, and security.
8. Develop end-to-end AI pipelines, integrating ML/DL solutions with real-world applications.

9. Implement MLOps best practices for model lifecycle management, versioning, monitoring, and retraining.
10. Collaborate with cross-functional teams to define AI-driven solutions that align with business objectives.
11. Stay updated with emerging AI/ML technologies, research advancements, and industry trends.
12. Mentor junior engineers and contribute to AI/ML knowledge sharing within the team.

Must Have:

1. Bachelor's or Master's degree in Computer Science, Engineering, Mathematics, or a related field.
2. 2+ years of experience in Python development, Machine Learning, and Deep Learning.
3. Strong programming skills in Python, with expertise in ML/DL frameworks such as TensorFlow, PyTorch, and Scikit-learn.
4. Experience with Time Series Forecasting techniques (ARIMA, LSTMs, Prophet, etc.).
5. Hands-on experience in Computer Vision (OpenCV, YOLO, Faster R-CNN, Vision Transformers, etc.).
6. Expertise in NLP and text processing (Transformer models, BERT, OCR technologies).
7. Experience developing RESTful APIs with FastAPI, Flask, or Django.
8. Proficiency in SQL and NoSQL databases (MySQL, PostgreSQL, MongoDB).
9. Experience in cloud-based model deployment on Azure, AWS, or GCP.
10. Strong knowledge of MLOps tools (MLflow, Kubeflow, SageMaker) for model management.
11. Familiarity with containerization (Docker, Kubernetes) for scalable AI deployments.
12. Strong understanding of software development best practices, including OOP, design patterns, and performance optimization.
13. Excellent problem-solving skills and ability to work in a fast-paced environment.
14. Strong communication and collaboration skills, with the ability to work in cross-functional teams.